

Homework One

Iacobescu Tudor, A6

27 oct 2019

1 Introduction

The following document details an experiment in which three optimization algorithms are ran against four multi-dimensional functions in order to try and obtain the functions' global minima.

The purposes of the experiment were initially to compare the three algorithms and how they operate, but it ended up being more interesting for the purposes of highlighting possible problems that pop up when implementing algorithms such as these. These problems don't make the algorithms unusable, but they do complicate their proper implementation.

The paper will describe the method of the experimental procedure, show the resulting data, and analyze it to try and find some interesting results – namely, what abnormalities are present in the data and what may have caused them.

2 Method

For the experiment, we will use two C++ implementations of Iterated Hill-climbing (with a First Ascent and Steepest Ascent approach respectively), as well as an implementation of Simulated Annealing. These algorithms will be tested against four common optimization algorithm testing functions, and their results will be compared.

2.1 Algorithm

All three of the algorithms rely on representing real numbers from a given interval $[a, b]$ as a string of bits. These bits simulate integers from 0 to 2^n (where n is the number of bits) corresponding roughly to values within the search interval. The number of bits necessary corresponds to the precision (10^{-p}) wanted:

$$n = \lceil \log_2 (b - a) \cdot 10^p \rceil$$

A point in d -dimensional space can be represented as a string of $d \cdot n$ bits. The set of a point's neighbors is that of all the bitstrings of Hamming distance of one from that point's bitstring.

```

00100110
10100110, 01100110, 00000110, 00110110
00101110, 00100010, 00100100, 00100111

```

Figure 1: A point represented by a string of 8 bits, and all of its neighbors.

Observation Assume the above is a point in one-dimensional space. Notice how neighbors whose differing bits are further to the left are farther in space from the original - exponentially farther, in fact. This property makes it so that neighbors end up being more common closer to the point in space than farther away, which effectively makes bigger jumps less likely.

Iterated Hillclimber This algorithm works by taking a number of randomly placed points in the domain of the function. For each candidate point, we then go through their neighbors to find a better candidate. We stop when we can no longer find a better neighbor. The result of the algorithm is the best point found by any of the chains above – the point found thus far where the function has the smallest value.

```

function iteratedHillclimber():
    vb = INITIAL_ASSUMED_BEST_CANDIDATE;
    for (i in 0..NUM_ATTEMPTS):
        vc = randomCandidate()
        repeat:
            vn = pickBetterNeighbor(vc)
            if (vc is null):
                break loop
        if (vc better than vb):
            vn = vc;
    return vb;

```

What separates Iterated Hillclimber into two variants is the way we pick the next neighbor to continue with for each candidate.

First ascent This method iterates through the neighbors only until a better one is found. The candidate is then replaced with said neighbor, and the cycle repeats.

```

function pickNeighbor(vc):
    neighbors = generateNeighbors(vc)
    for (vn in neighbors):
        if (vn better than vc)
            return vn
    return null

```

Steepest ascent This method iterates through all of a candidate's neighbors, and selects the best one (i.e. the one where the function takes the lowest value in) as the one to continue with.

```

function pickNeighbor(vc):
    neighbors = generateNeighbors(vc)
    vb = neighbors.pickAny()
    for (vn in neighbors):
        if (vn better than vb)
            vb = vn
    if (vb better than vc)
        return vb
    else
        return null

```

Simulated annealing This algorithm shares some similarities with Iterated Hillclimbing. Simulated Annealing takes inspiration from the real life process of annealing – a heat treatment process where a material such as metal or glass is repeatedly heated and cooled in order to make it more workable.

The algorithm works by picking a starting "temperature", and a random starting candidate point. A loop then goes through the candidate's neighbors – several methods can be used here, but I personally opted for randomly picking one a few times – and replacing the candidate upon finding a better neighbor. The thing that sets SA apart from Iterated Hillclimbing is the fact that even if a better neighbor is not found, there is still a chance that it might replace the candidate. This allows Simulated Annealing to go uphill as well instead of only downhill, which can allow it to have better odds of finding the global minimum instead of stopping at one of the local minima of the function, like Hillclimbing.

The chance that a worse neighbor will be picked to continue the algorithm is based on two factors - the difference between the current candidate and the neighbor (the worse the neighbor, the lower the chance) as well as the temperature of the current annealing stage (the higher the temperature, the higher the chance). Once a better neighbor is not found and a worse one fails this check as well, the temperature is lowered, and the cycle repeats until a certain stopping threshold.

The following is a simplification of the implementation of Simulated Annealing used for this experiment.

```
function simulatedAnnealing():
    T = ANNEALING_INITIAL_TEMPERATURE
    t = 0;
    vc = randomCandidate()
    while (T > ANNEALING_STOPPING_TRESHHOLD):
        consecutiveFailures = 0
        while (consecutiveFailures < ANNEALING_CONSEC_FAIL_LIM):
            vn = randomNeighbor(vc)

            if (vn better than vc):
                vc = vn
                consecutiveFailures = 0
            else:
                if (random0to1() < chanceTreshhold(T, vn, vc)):
                    vc = vn
                    consecutiveFailures = 0
                else:
                    consecutiveFailures++
        T = lowerTemp(T, t)
        t++

    return vc

function chanceTreshhold(T, vn, vc):
    diff = absolute difference between vn and vc's function values
    base = -(diff / T)
    return exp(base)

function lowerTemp(T, t):
    fraction = (t + 1) / (t + 2)
    return T * fraction
```

2.2 Functions

The four functions used for this experiment are the Sphere function, the Dixon & Price function, Michalewicz's function, and Rastrigin's function. The following subsections include the function expression and graph for each of these. The graph is for the two-dimensional version of each function, in its given search interval, with the vertical axis being the value of the function.

2.2.1 Sphere function

$$f(x_1 \cdots x_n) = \sum_{i=1}^n x_i^2, x_i \in [-5.12, 5.12]$$




Figure 2: Sphere function graph

2.2.2 Dixon & Price function

$$f(x_1 \cdots x_n) = (x_1 - 1)^2 + \sum_{i=2}^n i(2x_i^2 - x_{i-1})^2, x_i \in [-10, 10]$$




Figure 3: Dixon & Price function graph

2.2.3 Michalewicz function

$$f(x_1 \cdots x_n) = -\sum_{i=1}^n \sin(x_i) \left[\frac{ix_i^2}{\pi} \right]^{2m}, m = 10, x \in [0, \pi]$$




Figure 4: Michalewicz function graph

2.2.4 Rastrigin function

$$f(x_1 \cdots x_n) = 10n + \sum_{i=1}^n (x_i^2 - 10\cos(2\pi x_i)),$$




Figure 5: Rastrigin function graph

3 Experiment

The above algorithms were implemented in C++. For each of the four functions (for 2, 5 and 30 dimensions) all three of the algorithms were ran 30 times in parallel, each run limited to a time just above 600000ms (10 minutes). The results for each of the runs were transcribed in output files, which were then processed with a simple R script into a large L^AT_EX table, which was then manually split and edited. The resulting tables are shown below.

4 Results

The following tables show a summary of the result data for each algorithm, on each function and number of dimensions respectively. They present the mean, maximum, minimum and standard deviation of the results and times.

4.1 Iterated Hillclimber, First Ascent

function	d	rMean	rMax	rMin	rSDev	tMean	tMax	tMin	tSDev
dixon&price	2	0.03	0.14	0.00	0.04	1928.13	2122	1678	115.35
dixon&price	5	185.17	538.53	8.87	157.89	9573.10	10001	8859	304.60
dixon&price	30	1039584.90	1378580.00	639433.00	223446.61	48142.93	49648	45998	1378.71
Michalewicz	2	-1.80	-1.78	-1.80	0.01	1655.00	1806	1442	97.53
Michalewicz	5	-3.07	-2.68	-3.64	0.29	7348.27	7756	6786	285.30
Michalewicz	30	-8.33	-7.42	-9.67	0.52	28802.57	29682	27402	774.11
Rastrigin	2	0.53	1.24	0.00	0.47	1888.13	2058	1692	106.33
Rastrigin	5	18.96	29.31	5.00	4.86	8878.93	9446	4327	909.57
Rastrigin	30	368.88	405.04	323.45	20.66	27435.30	28453	26222	946.44
Sphere	2	0.00	0.00	0.00	0.00	1845.03	2043	1411	172.58
Sphere	5	1.56	3.90	0.18	0.82	8813.83	9245	8213	299.52
Sphere	30	133.49	151.02	117.92	8.90	27035.33	28100	25802	926.80

Table 1: IHC/FA result (r) and run time (t) statistics. Times are in ms.

4.2 Iterated Hillclimber, Steepest Ascent

function	d	rMean	rMax	rMin	rSDev	tMean	tMax	tMin	tSDev
dixon&price	2	0.01	0.05	0.00	0.01	6672.93	7034	5966	272.59
dixon&price	5	4.49	23.47	0.67	6.06	40694.50	42034	35141	1541.41
dixon&price	30	673153.87	930960.00	408164.00	134068.23	1017064.57	1029225	986768	9810.97
michalewicz	2	-1.80	-1.79	-1.80	0.00	5771.40	6124	5359	234.65
michalewicz	5	-3.73	-2.96	-4.64	0.41	34603.47	35394	33450	594.42
michalewicz	30	-8.86	-7.98	-9.82	0.44	887536.93	897716	870749	7719.99
rastrigin	2	0.36	1.15	0.00	0.38	6406.43	6777	5973	256.04
rastrigin	5	8.37	16.85	2.27	4.03	38223.17	39603	36586	1221.48
rastrigin	30	349.83	379.18	283.29	20.51	947533.87	960025	907167	10534.68
sphere	2	0.00	0.00	0.00	0.00	6381.27	6715	6008	251.18
sphere	5	0.00	0.02	0.00	0.00	39890.83	41305	33001	1655.24
sphere	30	111.89	129.99	84.18	11.71	938650.63	953482	901389	13566.75

Table 2: IHC/SA result (r) and run time (t) statistics. Times are in ms.

4.3 Simulated Annealing

function	d	rMean	rMax	rMin	rSDev	tMean	tMax	tMin	tSDev
dixon&price	2	0.99	5.98	0.04	1.35	5248.50	5667	4028	411.44
dixon&price	5	114510.50	289599.00	1731.43	75901.12	6172.77	7478	4917	595.88
dixon&price	30	3592796.33	6171060.00	1149410.00	1219380.78	26734.00	31553	19473	3138.98
michalewicz	2	-0.34	-0.00	-0.99	0.39	600005.90	600027	600001	7.95
michalewicz	5	-0.60	-0.01	-1.59	0.49	600005.53	600027	600001	7.34
michalewicz	30	-3.58	-1.46	-8.21	1.55	600003.73	600017	600001	4.49
rastrigin	2	2.04	6.10	0.17	1.17	50906.13	75110	21966	14973.56
rastrigin	5	59.77	84.04	32.02	12.23	207850.47	342680	99947	61872.17
rastrigin	30	563.78	672.59	384.66	53.07	558424.10	600028	257813	89574.33
sphere	2	0.92	4.37	0.04	0.91	270970.37	380789	93478	73123.96
sphere	5	43.26	82.65	12.80	17.59	582070.63	600016	409363	48450.30
sphere	30	269.46	359.97	160.07	44.08	600003.43	600021	600001	4.95

Table 3: Simulated Annealing result (r) and run time (t) statistics. Times are in ms.

5 Result analysis

Several interesting trends can be noticed in the results of the experimental data. These trends turn out to be harder to explain than first believed.

Dixon & Price inaccuracy One unexpected result of the experiment was that, by far, the least accurate results were obtained for the Dixon & Price function.

The algorithms have little trouble with it for two dimensions, but results rapidly deteriorate as the number of dimensions is increased. The obvious first possible cause for this would be the function's larger search interval ($[-10, 10]$) which due to the way that the search interval is mapped to bitstrings results in somewhat longer bitstrings overall, which increases the number of neighbors of any point.

Another possible factor is the shape of the D&P function's graph - many of the points yield very similar values, in the vicinity of the global minimum. Unfortunately, a definite correlation between either of these points and the experimental result could not be found.

Michalewicz annealing issues Another interesting experimental result is the fact that the Simulated Annealing algorithm hit its time limit in every run for the Michalewicz function. The results are consistently poor, and seem to be located on the function's plateau. On this plateau, results deviate very little from a value of zero. This is likely because the "trenches" that criss-cross the graph of Michalewicz are not very deep in absolute value, which boosts the chance that the algorithm will jump back onto the plateau. Because of this, the consecutive failure condition is very hard to meet, and the temperature tends to remain near its initial value, further boosting the odds that a worse neighbor will be accepted.

One interesting observation is that the results actually improve slightly for the 30 dimension runs. This may be due to the fact that the depth of the valleys in the function's graph increases with the number of dimensions (and so does the global minimum).

This problem highlights Simulated Annealing's greatest flaw: its efficiency is strongly influenced by the parameters chosen for the run, as well as subtleties in its implementation. It's almost certain that with slightly different parameters this problem would not have appeared, but it's likely another would have cropped up.

Iterated Hillclimber, Steepest Ascent time issues The 30-dimensional tests for all four of the functions with this method consistently took longer than the 600000ms time limit designed into the algorithm's implementation. The way the time limit is implemented is by breaking each individual attempt after it passes its allocated chunk of the time limit - for 1000 attempts, this means 6000ms for each one.

The time check is after the evaluation of the neighbors, so it's possible that the cause of this is that for 30 dimensions, the number of bits in a point's bitstring is so large, the evaluation of all neighbors ends up eating a significant chunk of time.

Iterated Hillclimber, First Ascent sphere inaccuracy It is bizarre that Hillclimber, FA should perform so poorly on the sphere function. The time limit is never reached, and yet the results are quite far from correct, a problem not present on the Steepest Ascent variant. I can't determine what caused this. I can't reproduce it in further experiments - even with 30 dimensions, First Ascent has little trouble with the sphere function – as it should, too, since the function graph has only one local/global minimum, which can easily be reached.

6 Conclusion

What was meant as an experiment to compare three simple optimization algorithms turned out to be an interesting trip into how the subtleties of the analyzed function can make or break the method used to analyze it. Seemingly innocuous things like domain size, the presence of plateaus, absolute value differences, or small changes with the number of dimensions used seem like they wouldn't make much of a difference, but end up causing otherwise functional algorithms to yield abnormal results.

Perhaps one way to look at all of the problems that plagued this experiment is a motivation for genetic algorithms to be used in the first place. After all, GAs are built around being able to handle the general case, and adapting their approach dynamically for each function encountered. It's quite possible that a genetic algorithm would not have had some of the problems present here.

Perhaps the next experiment will tell.

7 References

- Functions and graphs: http://www-optima.amp.i.kyoto-u.ac.jp/member/student/hedar/Hedar_files/TestGO_files/Page364.htm
- Test method information: <https://profs.info.uaic.ro/~pmihaela/GA/laborator2.html>